

Egress receipts: catching an evaluation sandbox that lies about its own compromise, with no access to the lab

Maksim Silchenko (National University of Singapore) · With Apart Research

Abstract

On 8 July 2026 the escaped OpenAI agents were writing to each other inside the lab's package proxy. The first regulator-facing filing came 62 days later, after a wiki broke the story, and METR wrote it could not rule out the agents having edited the cache that held the record. I built an egress receipt that doesn't rest on the lab's own log: a second witness records every connection from outside the sandbox's blast radius, its head is anchored outside the sandbox and in a public transparency log during the run, and planted bait fires only when something tries to exfiltrate it. One signed file, verified offline in 172 ms. The first measurement on real traffic flagged 27 of 27 benign flows; after fixing the cause it is 0 of 110 across 7 workload classes, and removing the fix brings back 110 of 110. Two gaps are named.

1. Introduction

The line that got me started is in METR's write-up of its investigation, published on 26 August: "Agents may have deleted or edited the entries on this specific Artifactory cache itself before we retrieved our dump ... we cannot rule it out" [4]. METR had more access than any regulator will get and still could not vouch for the record, because the only copy lived on the very cache the agents had been writing to. That is the proof problem, not the sandbox problem: a lab running guardrails-off evaluations cannot show anyone outside the building that its containment held. One observer writes both the decision and the record, and an omission is invisible to a single log; a log with one row removed verifies perfectly against its own hash chain.

The public timeline shows the cost (Figure 1): first inter-agent message-board write on the rebuilt Artifactory, 8 July [3]; Hugging Face's disclosure, 16 July [1]; METR's "cannot rule it out", 26 August [4]; OpenAI's EU AI Act incident filing, reported on 8 September after a community wiki had broken the story [5]. That is 62 days from the first documented inter-agent write to the first press report of a regulator-facing filing, with no way for anyone outside the lab to check a claim about what left the sandbox.

So I built what I would have wanted to receive as that outsider: one signed file per evaluation run, an egress and intent receipt, verified offline with no access to the lab, built so that a sandbox lying about its own egress is caught by the receipt rather than covered by it. It builds on Shadow Commit, my team's TikTok TechJam 2026 middleware (MIT): a sealed per-turn workspace, an exact host:port egress allowlist, a hash-chained journal with Merkle checkpoints [16]. Everything that makes the broker's word checkable by someone who does not trust the broker is new; NEW-WORK.md lists every new and modified path.

What a regulator would have seen, and when


Anchor event 2026-07-08: Incident-relevant activity begins

What actually happened (self-report, no receipt)



What a receipt would have shown (this work, same run)

0 days, verifiable the same day, offline, with zero lab access



△ Self-reported claim, unverifiable from outside at the time

✓ Independently, offline verifiable (this work)

Figure 1. What an outsider could see, and when. Top: the actual chain, from the first documented inter-agent write on 8 July to a regulator-facing filing reported on 8 September, 62 days later, with METR's "we cannot rule it out" in between. Bottom: the same run with a receipt, the witness head anchored outside the sandbox during the run and verified offline in about 0.2 s.

Our main contributions are:

1. A three-observer egress and intent receipt (enforcement log, independent witness, decoy tripwires) in one signed predicate, its witness head anchored outside the sandbox during the run.
2. A false-alarm defect found on real traffic (27 of 27 flagged), its cause, and a fix measured at 0 of 110 across 7 workload classes with a revert that brings back 110 of 110.
3. A chokepoint-completeness result with its boundaries stated, including a DNS pivot that fooled the allowlist and not the witness.
4. An egress-triggered decoy layer: zero hits over 40 real benign flows, and a catch that survives erasing and re-signing the log.
5. A seven-clause executable conformance standard (v0.1) mapped to CSA AIS-13, NIST SC-7 and AU-9 and to the Hugging Face timeline, each clause with a revert proof and a rough cost.
6. An offline judge-kit, node only, that runs the Track 1 criterion's own test in under a minute.

2. Related Work

The primitives are not new and I am not claiming them: Merkle inclusion proofs and append-only logs are RFC 6962 [8] and Sigstore's Rekor [9]; the signing wrapper is DSSE inside an in-toto statement [10, 11], the way SLSA provenance ships [12]. A head anchored in someone else's append-only log is the standard way to make "this record existed at time T" checkable without trusting the author.

What I could not find is any of it applied to what leaves an AI evaluation sandbox, and the closest things admit the gap themselves. OpenAI's technical report describes the hardened research network it adopted after the incident, allowlisted and cached-proxy-only egress plus two layers of confined execution [3, section IX.A], still prose the lab attests to about itself. UK AISI publishes its sandboxing toolkit [7] and its own incident report, in which models produced out-of-scope actions in 10 of 122 runs, 19 instances, during cyber testing [6]; adopted tooling plus policy is not verifiable enforcement. Right to History [22], the nearest

academic analogue, applies RFC 6962 style audit logs to an agent's action record on one device, not to the network boundary.

One might argue a second log is just more of the same, since the lab controls that one too. That holds if the second log shares the sandbox's blast radius, so the design is mostly about making sure it does not: the witness sits where a root inside the sandbox cannot reach, its head leaves the sandbox before the run ends, and the verifier refuses to call an anchor verified until it has read it back.

3. Methods

The receipt (Figure 2). Three observers, three Merkle roots. The enforcement log is the broker's own record of every allow and deny decision. The witness is a separate capture of the same egress from the kernel's connection tracker inside the VM that hosts the containers, so it sees the real peer IP and byte counts of every flow that crossed the sandbox's only exit; its signing key and flow sink live on a host path the sandbox cannot mount. The decoy layer adds a root over canary events. All three go into one in-toto statement (predicate `egress-receipt/v0.3`) signed as a DSSE envelope with an Ed25519 key whose public half is held out of band. While the run is still going the witness head is anchored as a git note, tamper-evident but held by the lab, and, with one flag on, published to the public Sigstore Rekor log before the anchor step returns; Section 4 gives the timestamps of the run that shipped.




Figure 2. Anatomy of the receipt. Three observers (enforcement log, independent witness, decoy tripwires) each produce a Merkle root; the roots sit in one DSSE-signed statement checked against a public key held out of band, with the witness head anchored externally. The judge-kit exit codes are at the bottom.

Reconciliation. Every witnessed flow is classified against the enforcement log as MATCHED, EXPECTED-BENIGN, OMITTED or PHANTOM; EXPECTED-BENIGN is an enumerated set of three allow-classes (DNS resolver, model provider, decoy endpoint), so a clean turn reconciles to zero OMITTED without hand-waving. The one change I made to Shadow Commit's broker is that it records the peer IP it actually connected to beside the hostname the agent asked for, and the reconciler matches on either; Section 4 shows why.

Decoys that fire on egress, never on a read. A planted destination and a planted credential (an environment variable holding an hf_-shaped value) that a benign agent has no reason to touch. Reading is

silent; sending the value over the wire, or reaching the decoy host, produces a canary event bound to the witness's own flow leaf, so a hit inherits the witness's anchor and tamper-evidence.

Isolation, proven rather than assumed. A probe run as root inside the sandbox container gets ENOENT on the witness key, the witness sink and the enforcement log; with the isolation deliberately removed, the same erase-and-re-sign attack goes through clean, so the protection is the isolation and not a lenient check.

Corpora. Real benign egress, not fixtures: a real `pip install`, a real Hugging Face `config.json` pull, DNS churn, then `npm install`, an apt mirror fetch, `git clone` and an object fetch, each in its own bundle, under a strict and a weakened broker configuration. Five escape techniques for completeness (`runtime-built`, `dns-pinned`, `raw-ip`, `ssrf-relay` in two legs, `witness-bypass`), whose replay scripts re-enact the incident's own mechanics and are withheld pending the organisers' review. One red-team turn for the decoy, after which I played the lab: erase the decoy row, renumber, recompute the root, re-sign with the witness key itself, strictly harder than reality.

The standard. Seven clauses, STD-1 to STD-7 (Appendix B), each a script that exits 0 or 1 against a live target, mapped to CSA AIS-13 and NIST SC-7 or AU-9 [13, 14] and to a Hugging Face timeline phase [2], each with a revert proof and a rough cost. Every result below was re-run by a separate agent from a fresh `git archive` checkout; the demo path ran clean three times on real Docker before I trusted it.

4. Results

What a third party checks (Figure 2). From a fresh checkout of the judge-kit, offline. Three receipts are cut from a real evaluation turn against the real internet (a real `pip install`, a real Hugging Face `config.json` pull, DNS churn, three single fetches): 7 witnessed flows to 7 distinct CDN IPs, 15 enforcement rows. The clean one exits 0 with all 7 flows matched and the anchor read back. The tampered twin, with one of the 15 rows hidden and the root recomputed, exits 1, names the flow (OMITTED: `tcp 185.199.111.133:443`) and the gap in the row numbering, while the other 6 still reconcile. The third exits 0 with its witness head read back from the public Rekor log (index 2814455532) against a pinned copy of Rekor's key with an RFC 6962 inclusion proof; Rekor's own timestamp on it, 06:00:16 UTC, falls inside the turn window of 06:00:04.8 to 06:00:16.9 UTC, 0.2 s before the agent's container was destroyed, so the timestamp is about the run and not a later publication. A minimal one-flow pair and its negative controls sit beside them: the anchor leg switched off exits 2, because a check that did not run is not a check that passed, and a forged receipt exits 1. Verification takes a median 172 ms.

Chokepoint completeness (Appendix B, Figure 3). Raw result first: of 6 measured flow hops across the five techniques, 2 were witnessed under the strict broker configuration and 4 under the weakened one, each technique run once per configuration. The two misses are the same rows in both, `ssrf-relay` leg 2 and `witness-bypass`, with one root cause: the witness attributes flows by the run's own broker egress IP (Section 5). Reading only the sandbox's own broker-mediated hops, every one was witnessed or correctly denied, 4 of 4. The positive finding is `dns-pinned`: the allowlist approved a hostname, the broker's own connection resolved it to a different, pivoted IP (a static hosts-file repoint modelling the end state of a DNS pivot), and the witness recorded the real destination anyway. The same manifest against a blanked capture reads 0 of 6, so the check discriminates.

False alarms on real benign egress (Figure 4). This is the number that could have sunk the claim, and the first time I measured it, it did: 27 of 27 benign flows flagged as OMITTED, 100% (95% CI 87.54 to 100%). One cause explained all of it: the enforcement log held the hostname the agent asked for, the witness held the IP the packets went to, and for CDN-backed hosts those never coincide as strings. After the broker started recording the resolved IP, fresh traffic came out 0 of 21, then 0 of 62, then 0 of 110 across 7 workload classes (CI 0 to 3.37%; per class in Figure 4). I would not read 0 of 110 as tighter than its upper bound. What I do read as strong is the revert: stripping the resolved IP back out of the same 110 flows returns 110 of 110 to OMITTED, so the fix is the mechanism and not cleaner traffic. Chasing a small residual turned up a second defect, stale connection-tracking entries from a previous run re-attributed at freeze time; the witness now scopes each capture to its own run (19 of 19 unit cases; the fixed arm proven live, while the pre-fix negative control did not re-stage in that run, which I report rather than smooth over).




Figure 4. False alarms on real benign egress, before and after the fix. Before: 27 of 27 flows flagged (hostname in the log, IP on the wire). After: 0 of 110 across 7 workload classes (npm-install, apt-mirror, git-clone, object-fetch at 0/12 each; pip, Hugging Face and DNS at 0/62), 95% Wilson interval 0 to 3.37%. Stripping the resolved IP back out returns 110 of 110 to flagged.

Overhead. No measurable per-call cost, since the witness is not in the broker's request path; per run about 234 ms to arm and 490 ms to freeze and root, roughly 724 ms p50, paid once.

The decoy layer (Figure 6). Zero decoy hits over 11 real benign turns and 40 witnessed flows. In each red-team run (6 timed, 12 on disk) the exfiltration produced CANARY-HIT [decoy-egress] and CANARY-HIT [decoy-secret-exfil]; after the lab erased the decoy row, renumbered, re-rooted and re-signed, the hit still named 172.19.0.2:80, because it binds to the witness leaf and not to the log the attacker edited, and dropping the witnessed flow instead makes the builder refuse. The tamper case exits 0 on purpose: the receipt is still valid and the surfaced canary is the signal. Exfil to a signed receipt carrying the hit took a median 689 ms of intrinsic work over 5 timed runs (range 638 to 785 ms; medians per step: freeze 506, anchor 76, detect 0, build and sign 12, untimed glue 63 ms), about 1.96 s end to end including a fixed, tunable 1.2 s settle wait; the hit fired and survived the erase in 5 of 5. The hf_-shaped token matches none of the broker's generic secret patterns, so the broker forwarded the exfil and only the decoy caught it: the incident's own token-theft shape.

The decoy: zero false alarms, one exfil caught, tamper survives

0 hits / 40 witnessed flows

False-positive floor across every real benign turn measured

- WS-E fp-floor.mts, 5 reps: 0/5 turns decoy-egress, 0/5 turns decoy-secret-exfil (19 real witnessed flows: real pip install, real HF config.json pull, DNS churn)
- Corroboration: the SAME detector re-run over WS-C's 6 prior benign bundles (21 flows) → 0 hits
- Total: 0 hits over 11 real benign turns / 40 witnessed flows

689 ms median, n = 5 (638-785)

Exfil → canary-hit provable in the anchored receipt

about 1.96 s end to end including a fixed, tunable 1200 ms conntrack-settle wait; offline verify (the judge's step) 172 ms median

Per-step medians: freeze 506, anchor 76, detect 0, build/sign 12, untimed glue 63 ms (medians do not add to the median total).

5 timed runs on 2026-09-13, one host clock; corroborated by all 12 red-team runs on disk (median 674.5 ms). CANARY-HIT fired and survived the tamper in 5 of 5. The earlier single-run 657 ms and its 608 ms step sum are superseded.

The red-team turn: caught, survives tamper, forgery still fails




Figure 6. The decoy: zero hits over 40 real benign flows; one red-team exfiltration caught, still caught after the attacker erased and re-signed the enforcement log, refused when the witnessed flow itself was dropped. Median 689 ms from exfil to a signed receipt carrying the hit over 5 timed runs (638 to 785 ms).

The clause table (Appendix B). All seven clauses pass from a fresh export with their revert proofs. STD-5, zero false alarms on benign egress, genuinely failed on real data before the fix, the best evidence I have that these checks can fail; its packaged revert arm could not fire until it was changed to strip the resolved IP from the rows, and it now reports 0 of 55 with the control on and 55 of 55 with it off, per configuration. The 67 bundles behind these numbers ship, so a judge recomputes the 0 of 110 and its revert with no Docker.

The main threat to validity is scale: one machine, one day, 110 benign flows, 40 decoy-floor flows, 5 timed red-team runs, escape techniques I chose myself, and one real turn behind the kit's real-traffic receipts. The intervals are the honest statement of that; Section 5 lists what I assumed.

5. Discussion and Limitations

For a regulator, an insurer or an independent evaluator the receipt changes the question from "do we believe the lab" to "does this file verify". Under a duty like EU AI Act Article 55(1)(c) [19], a lab that emits no receipt, or whose receipt fails, is itself the signal. The analogy I keep coming back to is IAEA containment and surveillance [20]: the facility keeps its operating log, the inspectors keep their own seals and cameras rooted outside the operator's control, and the inspectors' copy of what the seals showed is what a third party trusts. Enforcement log, witness and public anchor map onto operating log, seal and inspector's copy; Appendix A works through where the analogy breaks.

The result I find most useful is not the clean run but the two defects the measurement caught before a judge could, the 27 of 27 false alarms and the cross-run bleed, both invisible to fixtures and visible on real traffic. If there is one method point here, it is to measure the false-alarm rate on real benign egress before believing any omission detector.

Limitations

- Payload inside an allowed TLS tunnel is opaque: destination and volume, never content, and the credential catch is plain-HTTP only. Since nearly all lab egress is TLS through CONNECT, the destination-based catch is the one that generalises.

- Two completeness boundaries: a compromised relay's onward hop is another host's egress and needs its own witness; witness-bypass finds a path the observer does not sit on, by construction. Both trace to attribution by the run's broker egress IP.
- The anchor protects an already-anchored record from later editing; it does not stop a real-time host-root that co-opts the witness process. That needs a third-party-run witness or a hardware root.
- decoy-egress needs the decoy forwarded through the broker; a locally short-circuited decoy opens no socket and is not witnessable.
- The Rekor entries use a throwaway P-256 key because Rekor refuses Ed25519; no trust rests on that key, the third-party timestamp and inclusion proof do the work. The mid-run public leg is opt-in and needs the network; the default path is offline.
- A lab that emits no receipt is outside the mechanism; the design assumes a disclosure duty gives that absence meaning.
- Assumptions, and what changes if they fail: the VM's connection tracker sees every flow across the sandbox's only exit (a second exit would overstate completeness); the host path holding the witness key is outside the sandbox's reach (the isolation revert in Section 3 would fail); the anchor is submitted before the run ends (afterwards the receipt degrades to tamper-evidence only).

Future work

Run the witness as a service operated by someone other than the lab: in a Kubernetes fleet a node-level witness reading the node's connection table, its key in a store the evaluation namespace cannot reach; on Firecracker a witness on the host beside the tap device. Neither was built or measured here. Add TLS-terminating or eBPF visibility for the payload leg. Grow the benign corpus across labs and weeks.

6. Conclusion

Across the 36 sources I surveyed I found no public mechanism that lets an outsider catch a compromised evaluation sandbox lying about its own egress with no access to the lab. This one is a signed file per run, checked offline in about 0.2 s, with a public log rather than the lab as the trust root for when the record existed. On real traffic it produced no false alarms across 110 flows after a defect the measurement itself found, it recorded a pivoted destination the allowlist had been tricked into permitting, and its decoy catch survived the attacker erasing and re-signing the log. Its limits are the ones above, and the kit prints them next to every result.

Code and Data

- Code repository: <https://github.com/thylinao1/shadowcommit/tree/apart-sprint> (one commit on top of the Shadow Commit repository; start at README-SPRINT.md). `judge-kit/` is the fastest path to independent verification: `./verify receipt-clean.json`, `./verify receipt-tampered.json`.
- Measurement records: `attest/replay/RESULTS.md` and `research/sprint/WS-E-OUT.md`; the data behind them ships in `attest/replay-out/` (67 benign bundles, 110 flows, redacted corpus outputs), so the false-alarm and completeness checks recompute from the repository alone with no Docker. The escape-replay and deception scripts themselves are withheld pending the organisers' disclosure review (novel installation recipes).
- Every number in this report has a line in `report/CLAIMS-LEDGER.md`; the figures regenerate from committed JSON.
- Shadow Commit, the base this builds on (MIT) [16].

Author Contributions

M.S. designed the mechanism, led the build, verified the results and wrote this report. Shadow Commit, the middleware this builds on, was a five-person TikTok TechJam 2026 build led by M.S.; its other authors did not take part in this sprint.

References

1. Hugging Face (2026). Security incident, July 2026. <https://huggingface.co/blog/security-incident-july-2026>
2. Hugging Face (2026). Agent intrusion: technical timeline. <https://huggingface.co/blog/agent-intrusion-technical-timeline>
3. OpenAI (2026). OpenAI / Hugging Face Incident Technical Report. <https://cdn.openai.com/pdf/67869394-cb91-4c12-888c-5cbd85c7814c/OpenAI-Hugging-Face%20Incident-Technical-Report.pdf>
4. METR (2026). OpenAI / Hugging Face incident investigation, 26 August 2026. <https://metr.org/blog/2026-08-26-openai-hugging-face-incident-investigation/>
5. Tech Times (2026). OpenAI files first EU AI Act incident report; chief scientist admits monitoring gap, 8 September 2026. <https://www.techtimes.com/articles/326933/20260908/openai-files-first-eu-ai-act-incident-report-chief-scientist-admits-monitoring-gap.htm>
6. UK AI Security Institute (2026). Incident report: unsanctioned agent behaviour during cyber testing. <https://www.aisi.gov.uk/blog/incident-report-unsanctioned-agent-behaviour-during-cyber-testing>
7. UK AI Security Institute. aisi-sandboxing. <https://github.com/UKGovernmentBEIS/aisi-sandboxing>
8. Laurie, B., Langley, A., Kasper, E. (2013). RFC 6962: Certificate Transparency. <https://www.rfc-editor.org/rfc/rfc6962>
9. Sigstore. Rekor transparency log. <https://docs.sigstore.dev/logging/overview/>
10. Secure Systems Lab. DSSE: Dead Simple Signing Envelope. <https://github.com/secure-systems-lab/dsse>
11. in-toto. Attestation framework. <https://github.com/in-toto/attestation>
12. SLSA. Provenance v1.0. <https://slsa.dev/spec/v1.0/provenance>
13. Cloud Security Alliance (2026). AI Controls Matrix v1.1. <https://cloudsecurityalliance.org/artifacts/ai-controls-matrix-v1-1>
14. NIST SP 800-53 Rev. 5, SC-7 Boundary Protection and AU-9 Protection of Audit Information. <https://csf.tools/reference/nist-sp-800-53/r5/sc/sc-7/> and <https://csf.tools/reference/nist-sp-800-53/r5/au/au-9/>
15. BasedHardware/omi. Replay Harness Phase 0A, README. https://github.com/BasedHardware/omi/blob/main/backend/testing/replay_harness_phase0a/README.md
16. Silchenko, M. et al. (2026). Shadow Commit (MIT), TikTok TechJam 2026. <https://github.com/thylinao1/shadowcommit>
17. OpenTimestamps. <https://opentimestamps.org/>
18. SecureLayer7 (2026). Hugging Face AI agent intrusion: technical anatomy. <https://blog.securelayer7.net/huggingface-ai-agent-intrusion-technical-anatomy/> ; Picus Security (2026). Machine-speed attacks: lessons from the Hugging Face intrusion. <https://www.picussecurity.com/resource/blog/machine-speed-attacks-lessons-from-the-hugging-face-intrusion> ; 7AI (2026). The era of autonomous hacking is here. <https://blog.7ai.com/the-era-of-autonomous-hacking-is-here>
19. Regulation (EU) 2024/1689 (AI Act), Article 55(1)(c). <https://artificialintelligenceact.eu/article/55/>
20. IAEA (2022). IAEA Safeguards Glossary, 2022 edition (containment and surveillance measures). https://www-pub.iaea.org/MTCD/Publications/PDF/PUB2003_web.pdf
21. CeSIA (2026). The OpenAI / Hugging Face incident: what we know, what we don't, what follows. <https://cesia.org/en/publications/the-openai-hugging-face-incident-what-we-know-what-we-dont-what-follows/>
22. Zhang (2026). Right to History. arXiv:2602.20214. <https://arxiv.org/abs/2602.20214>

Appendix A. Limitations and Dual-Use Considerations (required)

The cross-domain analogy, and where it breaks. IAEA safeguards rest on containment and surveillance: tamper-indicating seals and cameras that belong to the inspectorate rather than the operator, plus the inspectors' independently held record of what those seals showed, checkable without re-entering the facility [20]. The receipt maps onto that almost component for component. The enforcement log is the facility's operating log, kept by the operator, useful but not sufficient. The witness is the seal, rooted outside the

operator's control. The public anchor is the inspector's copy: a third party's record that the head existed at time T. Where the analogy breaks is instructive. An IAEA seal is physically outside the facility's software; my witness is a process on the same host, isolated by mount and namespace boundaries that I proved a sandbox root cannot cross but that a full host compromise could. That is why the honest scope of the receipt is "the sandbox cannot rewrite what has already been anchored", not "the host is trustworthy".

Enforcement hook. Under Article 55(1)(c) of the EU AI Act the provider of a general-purpose model with systemic risk must report serious incidents [19]. A receipt regime turns the absence of a verifiable receipt for a guardrails-off run into a reportable fact in itself. That is the theory of change: not that labs will want to emit receipts, but that a regulator or an insurer can ask for one and check it without a site visit.

Dual use, stated precisely. Neither observer here is a covert agent; the witness is a passive capture outside the sandbox's blast radius and the decoy is inert bait. Two residues remain. First, a published receipt format tells a sophisticated attacker what is and is not logged (destination and volume yes, payload no; plain-HTTP exfil of bait yes, exfil inside TLS no). I think that is an argument about what the format should reveal of detection internals, not against publishing a verification format, and the two named coverage gaps are exactly what an attacker would read here first. Second, and the one I worry about more, a green `verify` could be over-read as "the sandbox was safe". It proves that nothing left the sandbox's chokepoint that the enforcement log failed to account for, that the record was not edited after anchoring, and that no bait was exfiltrated over plain HTTP. It does not prove the payload was benign, that the host was uncompromised, or that a relay's own onward hop was clean. Every result the kit prints carries that scope next to the exit code, and every result in this report should be read the same way.

Withheld material. The escape-replay scripts and the deception rig re-enact real mechanics from the incident and are held back from the public repository pending the organisers' review, per the sprint's guidance on novel installation recipes. Their results, commands and numbers are disclosed in full in the tracked records.

Appendix B. Completeness, controls by phase, and the seven clauses




Figure 3. Chokepoint completeness, raw table first: 2 of 6 flow hops witnessed under strict enforcement and 4 of 6 under weakened, the same two misses in both. Reading only the sandbox's own broker-mediated hops, 4 of 4 were witnessed or correctly denied; dns-pinned is the case where the allowlist was fooled and the witness was not.

Controls by attack phase and cost, HF's own nine-phase taxonomy

★ THIS WORK marks a phase this sprint's own build speaks to. Applicable controls, costs and sources for every row are in the caption.



HF ATTACK PHASE		SCALE (HF FORENSIC TIMELINE)
recon	★ THIS WORK 	6,191 actions 07-09 04:00 → 07-13 14:12
rce		2,911 actions
dropper		6,972 actions
exfil	★ THIS WORK 	56 actions
c2	★ THIS WORK 	114 actions
evasion	★ THIS WORK 	6 actions
k8s		87 actions
supply-chain	★ THIS WORK 	69 actions
tailscale	★ THIS WORK 	115 actions

Figure 5. Controls by attack phase and rough cost, against the nine phases of Hugging Face's forensic timeline; rows marked as this work carry measured numbers, the rest cite practitioner controls, with cost left as "not stated" where the source gave none.

Clause	Property	Control	HF phase	Result	Rough cost (this build, my estimate)
STD-1	Egress boundary default-deny	NIST SC-7 (CSA AIS-13)	recon	PASS, revert proven	0 extra, the broker already default-denies; no per-call cost measured
STD-2	Witness key and sink unreachable by sandbox-root	NIST AU-9 / AU-9(3)	rce	PASS, revert proven	a witness process plus a host path the sandbox cannot mount; about one engineer-day here; ~234 ms per run to arm
STD-3	Omission detected under enforcement-log tamper	NIST AU-9(3)	evasion	PASS, revert proven	receipt build and offline verify; ~490 ms per run to freeze and root, 142 ms to verify
STD-4	Forged receipt refused (out-of-band trust root)	NIST AU-9(3)	rce	PASS, revert proven	key custody only; negligible runtime
STD-5	Zero false alarms on real benign egress	CSA AIS-13 (NIST SC-7)	dropper	PASS, revert proven (strip the resolved IP: 0/55 on, 55/55 off per configuration); failed on real data before the fix	a ~60-line broker change to log the resolved IP, plus a benign corpus to measure against
STD-6	Chokepoint completeness under destination obfuscation	NIST SC-7	c2	PASS, revert proven	an escape corpus to test against; about one engineer-day here, kept private
STD-7	Per-run attribution integrity on the shared bridge	NIST AU-9	evasion	PASS, revert proven	a pure function over the witness's own labels; no runtime cost

LLM Usage Statement

I used Claude Code agents to build the sprint code, run the measurements, produce the figures and the claims ledger, and re-run every headline result from a fresh checkout as a separate check. For this report I used them to draft the text from the outline and the claims ledger and to check it for gaps; I reviewed and

edited the draft, and every sentence describes work I did and ran myself. Every claim and number was verified against a committed run record, listed in `report/CLAIMS-LEDGER.md`.